

# Methods and Tools

---

# Interactive Lectures

**All lectures in the course will be **interactive****

They contain running code, as well as theory!

- Presented and discussed in frontal lectures...
- ...You can download PDFs
- ...But you will also be able to **make changes and experiment**

**From a software perspective, the workhorses of this approach are:**

- Jupyter notebooks for the presentation & interaction
- Docker containers for the setup and distribution

**Both are widely used systems:**


- Jupyter is a user favorite when it comes to data science
- Docker is a state-of-the-art system for managing services

# A Few Words about Docker

---

# Docker

Docker is a system for running software in "containers"




Think of a container as a **lightweight virtual machine**:

- (Essentially) the same level of isolation
- ...But smaller disk footprint, faster setup and operation, etc.

# Docker

Docker is a system for running software in "containers"



Using containers has many advantages:

- Multiple environments on the same machine
- Improved isolation, robustness, and reproducibility
- Easier replication
- Scalability of cloud services...

# Docker

**During this course we will see many problems**

...And tackle them with **many techniques**:

- Classical Machine Learning
- Deep Learning
- Statistics
- Signal processing
- Declarative optimization
- Differential Equations
- ...

Managing dependences can become hellish

# Docker

**With docker, we can simply use a different container per case study**

Inside each container we will have:

- All the needed libraries & tools
- A running instance of a Jupyter server

In the host machine (your PC):

- We will just open a browser...
- ...And connect to the Jupyter server

## Two key concepts in Docker

- A container is a (sort of) running, lightweight, Virtual Machine
- An image is (sort of) the content of the hard disk of the VM

The image can be used to instantiate multiple containers

# Building an Image

## Images in docker are built by:

- Starting from a base image on [Docker Hub](#)
- Copying content between the host and the container
- Running commands in the container

## The process is controlled via a Dockerfile

- Just a text file with a specific syntax
- There is an [extensive reference](#), but we only care about a few commands

## To build an image, we can use:

```
docker build .
```

...From the directory with the Dockerfile



# An Example Dockerfile

This is a simple **Dockerfile** example:

```
FROM python:3.8
RUN pip install jupyter pandas sklearn matplotlib ipyml RISE
COPY . /app
WORKDIR /app/notebooks
CMD ["jupyter", "notebook", "--port=8888", "--no-browser", "--ip=0.0.0.0", "--allow-root"]
```

- The FROM keyword specifies the base image

# An Example Dockerfile

This is a simple **Dockerfile** example:

```
FROM python:3.8
RUN pip install jupyter pandas sklearn matplotlib ipyml RISE
COPY . /app
WORKDIR /app/notebooks
CMD ["jupyter", "notebook", "--port=8888", "--no-browser", "--ip=0.0.0.0", "--allow-root"]
```

- The RUN keyword runs a command
- In our case, we install a number of python packages

# An Example Dockerfile

This is a simple **Dockerfile** example:

```
FROM python:3.8
RUN pip install jupyter pandas sklearn matplotlib ipyml RISE
COPY . /app
WORKDIR /app/notebooks
CMD ["jupyter", "notebook", "--port=8888", "--no-browser", "--ip=0.0.0.0", "--allow-root"]
```

- The COPY keyword transfers data from the host to the container
- The first path refers to the host
- The second path to the container

# An Example Dockerfile

This is a simple **Dockerfile** example:

```
FROM python:3.8
RUN pip install jupyter pandas sklearn matplotlib ipyml RISE
COPY . /app
WORKDIR /app/notebooks
CMD ["jupyter", "notebook", "--port=8888", "--no-browser", "--ip=0.0.0.0", "--allow-root"]
```

- The WORKDIR changes the current directory in the container
- It's like running `cd` in the container

# An Example Dockerfile

This is a simple **Dockerfile** example:

```
FROM python:3.8
RUN pip install jupyter pandas sklearn matplotlib ipyml RISE
COPY . /app
WORKDIR /app/notebooks
CMD ["jupyter", "notebook", "--port=8888", "--no-browser", "--ip=0.0.0.0", "--allow-root"]
```

- The CMD keyword is triggered only when we **run** a container
- It's the first command that the container should execute
- It does nothing when building an image

# An Example Dockerfile

**When we run `docker build .` for our file:**

- The docker daemon downloads the base image, if not already available
- A container is started
- All operations in the Dockerfile are executed
- The resulting container is dumped, to create an **image**

**You can check that a new image has been built using:**

```
docker image ls
```

You will see an entry with no name:

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
<none>	<none>	96b910c1514f	3 seconds ago	1.36GB

# An Example Dockerfile

**You can remove an image with:**

```
docker image rm <image name or id>
```

- Useful to free space, however...
- ...Images are incremental! Docker stores only the differences
- ...So, don't worry too much about space usage

**You can remove all images with no running container with:**

```
docker image prune
```

# Running a Container

**You can instantiate and run a container with:**

```
docker run <image name or id>
```

- The container `stdout` will be piped (i.e. connected) to your terminal
- By default, this is not the case for `stdin`
- You can make the container interactive with the `-it` options
- You can autoremove the container at the end with `--rm`
- You can sync folders in the host and on the container using volumes

The documentation is extensive



# Running a Container

You can obtain the list of all containers with:

```
docker ps
```

- The option `-a` shows all containers (incl. those that are stopped)

You can remove a container with

```
docker rm <container id>
```

- As you see, it's a very flexible system
- ...But also a bit complex

**That's why we will automate most operations using Docker Compose**

# Docker Compose

**Docker Compose** is a tool to help the management of containers

In a second `docker-compose.yml` file, you specify:

- Which "services" (i.e. container) should be built and run
- How to build them
- Which options to use when running them
- ...

All in a human-readable, declarative format

# A Docker Compose Example

Let's see a simple `docker-compose.yml` for one lecture:

```
version: '2.0'
services:
  jupyter:
    build: .
    ports:
      - "8888:8888"
    volumes:
      - .:/app
```

- `version` refers to the Docker Compose syntax
- `services` is followed by a list of the containers
- `jupyter` is our service
- `build` specifies where the Dockerfile can be found

# A Docker Compose Example

Let's see a simple `docker-compose.yml` for one lecture:

```
version: '2.0'
services:
  jupyter:
    build: .
    ports:
      - "8888:8888"
    volumes:
      - .:/app
```

- `ports` tells which ports to expose to docker `run`
- `volumes` specifies which folders to sync
- In our case `"/app/notebooks"` on the container
- ...Will actually be `"/notebooks"` on the host

# Benefits of Using Docker Compose

**We need to use one more tool, but now we can:**

Build and run a container with:

```
docker compose up
```

- The command can also restart a stopped container

Stop the container with CTRL+C, or with:

```
docker compose stop
```

Stop and remove the container with:

```
docker compose down
```

...Which is considerably simpler than before!

# Our Jupyter Setup

---

# Our Setup

## We will often work with this development setup

The folder with the notebooks is structured as follows:

```
notebook1.pynb
notebook2.pynb
...
util      <-- module
assets    <-- images and such
rise.css  <-- for the "slide" mode
```

# Our Setup

## We will often work with this development setup

The folder with the notebooks is structured as follows:

```
notebook1.pynb
notebook2.pynb
...
util +-- __init__.py
      +-- XYZ.py      <-- submodule
      +-- YZX.py      <-- submodule
      +-- ...
font
rise.css
```

**The most important part:** we'll use **modules** besides notebooks



# Our Setup

## Working with modules provides some advantages:

We do not need to keep all our code in the notebooks. We can:

- Share functions **between cells**
- Share functions **between notebooks**
- IDEs can offer **more functionality** if they recognize a module

## ...But also a significant disadvantage:

- Python modules are compiled first when loaded...
- ...The loaded version is **not updated** when the source changes

This is very inconvenient at development time

# Our Setup

**We can circumvent this thanks to Jupyter "magic" extensions**

The first one is the "autoreload" extension

```
In [1]: %load_ext autoreload  
        %autoreload 2
```

- `load_ext` will enable the extension
- `autoreload 2` will reload all modules before code execution

**This is *inefficient, but convenient* during development**

- Together with the use of volumes (in docker-compose)...
- ...This allows us to update the code without re-building the docker image

# Starting a Notebook

Let's look back to the `CMD` keyword in our Dockerfile:

```
CMD ["jupyter", "notebook", "--port=8888", "--no-browser", \  
      "--ip=0.0.0.0", "--allow-root"]
```

This is translated to:

```
jupyter notebook --port=8888 --no-browser --ip=0.0.0.0 --allow-root
```

- `--port 8888`: the server listen on port 8888
- `--no-browser`: do not open the browser (there's no browser in the container)
- `--ip=0.0.0.0`: listen on all network interfaces
- `--allow-root`: we operate as `root` (admin) on the container

# Starting a Notebook

When we run:

```
docker compose up
```

The output will look like:

```
Starting ad_stat_jupyter_1 ... done
...
...Use Control-C to stop this server and shut down all kernels...
...
...To access the notebook...
...copy and paste one of these URLs:
...   http://34b908cf2362:8888/?token=82e337a2be9915cdebce276bf...
... or http://127.0.0.1:8888/?token=82e337a2be9915cdebce276bf...
```

- The last URL can be **copy-pasted in your favorite browser**

# Starting a Notebook

When we run:

```
docker compose up
```

The output will look like:

```
Starting ad_stat_jupyter_1 ... done
...
...Use Control-C to stop this server and shut down all kernels...
...
...To access the notebook...
...copy and paste one of these URLs:
...   http://34b908cf2362:8888/?token=82e337a2be9915cdebce276bf...
... or http://127.0.0.1:8888/?token=82e337a2be9915cdebce276bf...
```

- The token is cached by the browser and grants access to the notebooks

# Our Setup

- We will sometimes use the `ipymp1` package and the widget jupyter magic
- This will display basic tools to rescale and zoom images

```
In [2]: #!/matplotlib widget  
from matplotlib import pyplot as plt  
import numpy as np  
x = np.linspace(0, 2*np.pi, 100)  
plt.figure(figsize=(9, 3))  
plt.plot(x, np.sin(x))  
plt.tight_layout()
```

