

# A Formalization for Anomaly Detection

---

**How do we formalize our problem?**

# Problem Formalization

**A possible approach: we characterize the data distribution**

If we can **estimate the probability** of every occurring observation  $x$

- We can choose a size for the car pool
- ...Then we can spot anomalies based on their **low probability**

After all, anomalies are **rare events**, by definition


**We turn a liability into a strenght!**

# Problem Formalization

**We can check our intuition on our data**

This is (roughly) the distribution **over all the data**

```
In [3]: vmax = data['value'].max()  
util.plot_histogram(data['value'], vmax=vmax, bins=20)
```




# Problem Formalization

## We can check our intuition on our data

This is (roughly) the distribution around the first anomaly:

```
In [4]: w0_start, w0_end = windows.loc[0]['begin'], windows.loc[0]['end']  
data_anomaly0 = data[(data.index >= w0_start) & (data.index < w0_end)]  
util.plot_histogram(data_anomaly0['value'], vmax=vmax, bins=30)
```



- It seems indeed that there's a significant difference

# Problem Formalization

**When we reach this stage, it's a good idea for **formalize** our problem**

We can characterize a continuous distribution via its **density**

- Given a random variable  $X$  with values  $x$
- ...We care about its **Probability Density Function**  $f(x)$

**Since anomalies are assumed to be unlikely**

...Our detection condition can be stated as:

$$f(x) \leq \varepsilon$$

- Where  $\varepsilon$  is a (scalar) threshold

# Problem Formalization

**When we reach this stage, it's a good idea for **formalize** our problem**

We can characterize a continuous distribution via its **density**

- Given a random variable  $X$  with values  $x$
- ...We care about its **Probability Density Function**  $f(x)$

**Since anomalies are assumed to be unlikely**

...Our detection condition can be stated as:

$$f(x) \leq \varepsilon$$

- Where  $\varepsilon$  is a (scalar) threshold

**What do we need to make this work?**

# Density Estimation

**We need one way to estimate probability densities**

For some random process with n-dimensional variable  $\mathbf{x}$ :

- Given the true density function  $f(\mathbf{x}) : \mathbb{R}^n \rightarrow \mathbb{R}^+$
- ...And a second function  $\hat{f}(\mathbf{x}, \theta)$  with the same input, and parameters  $\theta$

We want to make the two as similar as possible



# Density Estimation

**We need one way to estimate probability densities**

For some random process with n-dimensional variable  $\mathbf{x}$ :

- Given the true density function  $f(\mathbf{x}) : \mathbb{R}^n \rightarrow \mathbb{R}^+$
- ...And a second function  $\hat{f}(\mathbf{x}, \theta)$  with the same input, and parameters  $\theta$

We want to make the two as similar as possible

**Can we obtain that using supervised learning?**

Given some suitable loss function  $L(y, \hat{y})$ , we would need to solve:

$$\operatorname{argmin}_{\theta} L(\hat{f}(\mathbf{x}, \theta), f(\mathbf{x}))$$

- where  $\mathbf{x}$  represents the training data

# Density Estimation

Unfortunately, this approach **cannot work**

...Because typically **we do not have access** to the true density  $f^*$

Density estimation is an **unsupervised** learning problem

It can be solved via a number of techniques:

- Simple histograms
- Kernel Density Estimation
- Gaussian Mixture Models
- Normalizing Flows
- Non Volume Preserving (NVP) transformations

Which one shall we pick?

# Our Friend, Occam's Razor

## We will go with Occam's razor

It's a philosophical principle stating that:

**Between two hypotheses, the simpler one is usually correct**



# Our Friend, Occam's Razor

## We will go with Occam's razor

It's a philosophical principle stating that:

**Between two hypotheses, the simpler one is usually correct**

**In practice, it's often a good idea to start with a simple approach**

- If it works well, then you have a solution
- If it does not, then you have a baseline

In both cases, you win!

# Our Friend, Occam's Razor

## We will go with Occam's razor

It's a philosophical principle stating that:

**Between two hypotheses, the simpler one is usually correct**

**In practice, it's often a good idea to start with a simple approach**

- If it works well, then you have a solution
- If it does not, then you have a baseline

In both cases, you win!

**For this example, we will pick Kernel Density Estimation**

# Kernel Density Estimation

---

# Kernel Density Estimation

In **Kernel Density Estimation (KDE)**, the main idea is that:

- Wherever (in input space) there is a sample
- ...It's likely that there are more

So, we assume that each training sample is the center for a density "kernel"

**Formally, each kernel  $K(x, h)$  is just a valid PDF:**


- $x$  is the input variable (scalar or vector)
- $h$  is a parameter (resp. scalar or matrix) called **bandwidth**

Typical kernels: Gaussian, exponential, cosine, linear...

# Kernels

An example with one sample and a **Guassian** kernel:

```
In [6]: x = np.array(0.5).reshape(1,1) # single sample
kde = KernelDensity(kernel='gaussian', bandwidth=0.1) # build the estimator
kde.fit(x) # fit the estimator on the data
# We use a plotting function from our module
util.plot_density_estimator_1D(kde, xr=np.linspace(0, 1, 200))
ymin, ymax = plt.ylim()
plt.vlines(x, ymin, ymax, color='tab:red')
plt.ylim((ymin, ymax)); # ; = suppress output
```






# Kernels

An example with one sample and a **Tophat** kernel:


```
In [7]: x = np.array(0.5).reshape(1,1) # single sample
kde = KernelDensity(kernel='tophat', bandwidth=0.1) # build the estimator
kde.fit(x) # fit the estimator on the data
# We use a plotting function from our module
util.plot_density_estimator_1D(kde, xr=np.linspace(0, 1, 200))
ymin, ymax = plt.ylim()
plt.vlines(x, ymin, ymax, color='tab:red')
plt.ylim((ymin, ymax)); # ; = suppress output
```



# Kernel

An example with one sample and a **linear** kernel:

```
In [8]: x = np.array(0.5).reshape(1,1) # single sample
kde = KernelDensity(kernel='linear', bandwidth=0.1) # build the estimator
kde.fit(x) # fit the estimator on the data
# We use a plotting function from our module
util.plot_density_estimator_1D(kde, xr=np.linspace(0, 1, 200))
ymin, ymax = plt.ylim()
plt.vlines(x, ymin, ymax, color='tab:red')
plt.ylim((ymin, ymax)); # ; = suppress output
```



# Kernels

As an example, a **Gaussian kernel** in sklearn is given by:

$$K(x, h) \propto e^{-\frac{x^2}{2h^2}}$$

- The  $\propto$  ("proportional to")

The function is similar to a the PDF of the Normal distribution:

- The mean can be interpreted as **0**
- $h$  plays the role of the standard deviation
- ...And scikit learn handles normalization

# Kernel Re-centering

Since the "mean" is 0, the kernel is **centered on 0**

All kernels in KDE are by default **zero-centered**

- ...But we need to place them over each sample
- How can this be done?

# Kernel Re-centering

Since the "mean" is 0, the kernel is **centered on 0**

All kernels in KDE are by default **zero-centered**

- ...But we need to place them over each sample
- How can this be done?

**We can use an **affine transformation** (like in the scale/location trick)**

In practice, the expression:

$$K(x - \mu, h)$$

- ...Gives the value of a kernel **centered on  $\mu$**
- ...Computed for the value  $x$

# Kernel Density Estimation

The estimated density of any point is obtained as a **kernel average**:

$$f(x, \bar{x}, h) = \frac{1}{m} \sum_{i=0}^m K(x - \bar{x}_i, h)$$

- $x$  is the input for which we want an estimate
- $\bar{x}_i$  is sequence of the  $m$  training samples
- $x - \bar{x}_i$  is the difference between  $x$  and the  $i$ -th training sample

**By changing the kernel function:**

- We can adjust the properties of the distribution (e.g. smoothness)
- Typically, the choice is based on prior domain knowledge

# Kernel Density Estimation

**KDE models are not trained in the usual sense**

...But they store internally **all the training samples**

- I.e. the training set is **part of the model parameters**
- This is a property common to most kernel models

# Kernel Density Estimation

**KDE models are not trained in the usual sense**

...But they store internally **all the training samples**

- I.e. the training set is **part of the model parameters**
- This is a property common to most kernel models

**There is one thing that we need to train, i.e. the bandwidth  $h$**

- We will see a general approach later in the course
- ...But in the **univariate** case we can apply a rule of thumb:

$$h = 0.9 \min \left( \hat{\sigma}, \frac{IQR}{1.34} \right) m^{-\frac{1}{5}}$$

Where ***IQR*** is the inter-quartile range



# Kernel Density Estimation

An example with two samples and a Gaussian kernel:

```
In [9]: x = np.array([0.25, 0.75]).reshape(-1,1) # two sample, univariate
kde = KernelDensity(kernel='gaussian', bandwidth=0.1) # build the estimator
kde.fit(x) # fit the estimator on the data
util.plot_density_estimator_1D(kde, xr=np.linspace(0, 1, 200))
ymin, ymax = plt.ylim()
plt.vlines(x, ymin, ymax, color='tab:red')
plt.ylim((ymin, ymax)); # ; = suppress output
```

